

Parameswaran Raman

CONTACT INFORMATION	115 Leonard Street, Apt A Santa Cruz, CA 95060	☎: (408) 306-4462, e-mail: params@ucsc.edu homepage: http://people.ucsc.edu/~praman1
RESEARCH INTERESTS	Optimization for Machine Learning, Large Scale Bayesian Inference, Learning to Rank	
EDUCATION	PhD - Computer Science , UC Santa Cruz Thesis Advisor: S.V.N Vishwanathan (Transferred from Purdue University (2013 - 2014))	Aug 2013 - present
	Masters - Computer Science , Georgia Institute of Technology	Aug 2009 - May 2011
	MSc (Integrated) - Software Engineering , PSG College of Technology, India	2003 - 2008
	<i>Graduate Research Assistant, UC Santa Cruz</i>	Aug 2013 - present
RESEARCH EXPERIENCE	<ul style="list-style-type: none">• [In Progress] Efficient methods for large-scale bayesian inference in mixture of exponential families.• [In Submission] Exploiting Double Separability for Scaling up Distributed Multinomial Logistic Regression problems to very large number of classes and examples.• [Accepted for NIPS 2014] Proposed a new Learning to Rank algorithm (RoBiRank) inspired by Robust Binary Classification, which directly bounds NDCG. Demonstrated scaling on large datasets.	
	<i>Research Intern - Adobe Research (STL)</i> (Mentors: Hung Bui, Branislav Kveton) Worked on clustering user-behavior in Adobe analytics data using user url click information.	Summer 2016
	<i>Research Intern - Microsoft (Cloud and Information Services Lab)</i> (Mentors: Keerthi Selvaraj, Dhruv Mahajan) Worked on the problem of extrapolating learning curves in machine learning. The goal was to study if it was feasible to use information from the models learnt on various sizes of small bites of data to extrapolate performance of the algorithm on the full data.	Summer 2015
	<i>Research Intern - Search Relevance (SNA), LinkedIn</i> (Mentors: Viet Ha-Thuc, Shakti Sinha) Worked on resolving issues of sample bias and position bias present in learning to rank systems with implicit feedback. Explored a variety of approaches, and ran offline and online experiments on LinkedIn Job Search portal to verify the effectiveness of the models.	Summer 2014
	<i>Graduate Research Assistant, ITAP - Purdue University</i> [Accepted for EDM 2014] Applied machine learning on educational mobile & web apps data to infer relevance of user posts to the lecture topics. Results helped improve student engagement in classrooms.	Aug 2013 - Aug 2014
	<i>Independent Research, Info Lab, Stanford University</i> Explored the scope of using learning methods in crowd-sourcing systems, to improve label complexity and quality of judgements among workers. In addition, experimentally analyzed the effects of using various interfaces for categorization of items in a taxonomy, simultaneously trying to model their error rates.	Summer 2012
	<i>Graduate Research Assistant, Sonification Lab, Georgia Tech</i> [Work appeared in several demos and accepted for ICAD 2010, CSUN 2010] Prototyped tools to demonstrate key ideas that came up in two projects Auditory Menus and In-Vehicle Assistive Technology (IVAT). Used machine learning for driver mood detection & providing alerts.	Aug 2009 - May 2011
	<i>Research Assistant, Dept of Maths & Computer Applications, PSG Tech</i>	2007 - 2008

[Accepted for DCCA Jordan 2007 and PETRA 2008] Worked with Dr Nadarajan and Dr Maytham Safar to propose effective cache replacement policies for Location-Dependent Data in mobile environments and implement tools for evaluating them.

INDUSTRY EXPERIENCE

Software Engineer, Yahoo!, Sunnyvale

Jul 2011 - Jul 2013

- Worked for the Personalization group, on an entity detection/resolution system used by all personalization services. Used machine learning and NLP to detect word/phrase boundaries and rank extracted entities. Built a Knowledge Graph from scratch to power Yahoo! search products.
- Worked on the Web of Objects project, to create a semantic knowledge base of entities to enable personalization. Designed features for Entity Matching models and wrote tools to evaluate them.
- Worked on Apache Oozie (a widely used job scheduler for Hadoop), implementing several features and fixing bugs in the system.

Software Engineering Intern, Intel, Chandler

Summer 2010

Developed a searching & indexing infrastructure to help silicon engineers find relevant product design information. Gathered requirements, developed & tested the system, and deployed in production.

Application Developer, ThoughtWorks, Bangalore

Jun 2008 - Jul 2009

Designed and implemented web-services for the train ticket retailing system - thetrainline.com. Worked in a fully agile setup following iterative test-driven software development.

PUBLICATIONS & POSTERS

- Parameswaran Raman, Jiong Zhang, Hsiang-Fu Yu, Shihao Ji, S.V.N Vishwanathan. “**Extreme Stochastic Variational Inference: Distributed and Asynchronous**,” *arXiv* pre-print.
- Parameswaran Raman, Shin Matsushima, Xinhua Zhang, Hyokun Yun, S.V.N Vishwanathan. “**DS-MLR: Exploiting Double Separability for Scaling up Distributed Multinomial Logistic Regression**,” *arXiv* pre-print.
- Hyokun Yun, Parameswaran Raman, S.V.N. Vishwanathan. “**Ranking via Robust Binary Classification and Parallel Parameter Estimation in Large-Scale Data**,” *NIPS*. 2014.
- Mariheida Córdova Sánchez, Parameswaran Raman, Luo Si, Jason Fish. “**Relevancy Prediction of Micro-blog Questions in an Educational Setting**,” in *Proceedings of the 7th International Conference on Educational Data Mining, EDM*. 2014.
- Myounghoon “Philart” Jeon, Benjamin Davison, Jeff Wilson, Parameswaran Raman, Bruce N. Walker. “**Reducing repetitive development tasks in auditory menu displays with the auditory menu library**,” in *Proceedings of the 16th International Conference on Auditory Display (ICAD)*. 2010.
- Parameswaran Raman, Narayanan Ramakrishnan, Manohar Ganesan, Gourab Kar, Dr Gregory D. Abowd. “**PiX-C: Express and Communicate (Augmenting Communication with Visual Input for Children in the Autism Spectrum)**,” in *Poster presented at the UIST Student Innovation Contest*. 2010.
- Mary Magdalene Jane, Parameswaran Raman, Maytham Safar, Nadarajan R. “**PINE-guided cache replacement policy for location-dependent data in mobile environment**,” in *Proceedings of the First international conference on Pervasive Technologies Related to Assistive Environments, PETRA*. 2008.

ACADEMIC SERVICES

- Reviewer for UAI 2014, AISTATS 2015, COLT 2015, JMLR 2015, AISTATS 2016, ICML 2016
- Developed and taught a bootcamp on Unix & Shell commands to undergrad students at UCSC
- TA for Machine Learning (Grad Level) course at UCSC

- HONORS/AWARDS
- Graduate Research Assistantship, UC Santa Cruz
 - Graduate Research Assistantship, Georgia Tech
 - Winner of Facebook Hackathon 2010 at Georgia Tech & finalists at FB HQ
 - Finalist for the poster presentation at UIST Student Innovation Contest 2010
- COMPUTER SKILLS
- C++, Java, Python, Matlab, MPI, Hadoop, Apache Spark, Unix, Lisp
- REFERENCES
- S.V.N Vishwanathan (vishy@ucsc.edu), Professor in Computer Science, UC Santa Cruz
 - Hung Bui (bui.h.hung@gmail.com), Senior Machine Learning Scientist, Adobe Research, San Jose
 - S. Sathiya Keerthi (keerthi@microsoft.com), Principal Scientist Manager, Microsoft (Cloud Information Services Lab), Mountain View
 - Viet Ha-Thuc (vhathuc@linkedin.com), Senior Data Scientist, LinkedIn, Mountain View
 - Reiner Kraft (reiner@yahoo-inc.com), Director of Engineering, Yahoo!, Sunnyvale